LAB 1

EXPERIMENTING WITH OPENFLOW IN MININET

AKHILA NAIR

01743358

Turn in Date: 2/12/18

Estimated Hours of Working: 7 Hours.

1) Purpose:

- Installing and configuring virtual box an SDN hub.
- Executing a packet flow network that can be traced graphically.
- Creating a virtual network with the switch, controller and hosts.
- Understanding basic mininet commands and its functions.
- Understanding Openflow protocol and their communication between switch and controller.
- Experimenting with Wireshark and analysing Loopback interfaces.

2) Setup:

Lab 1 procedure has following steps:

2.1: Setting up Virtual Machine –

- Open Virtual Box and import SDN Hub.
- IP Address of virtual machine can be found with commands – sudo dhclient eth1 and ifconfig.

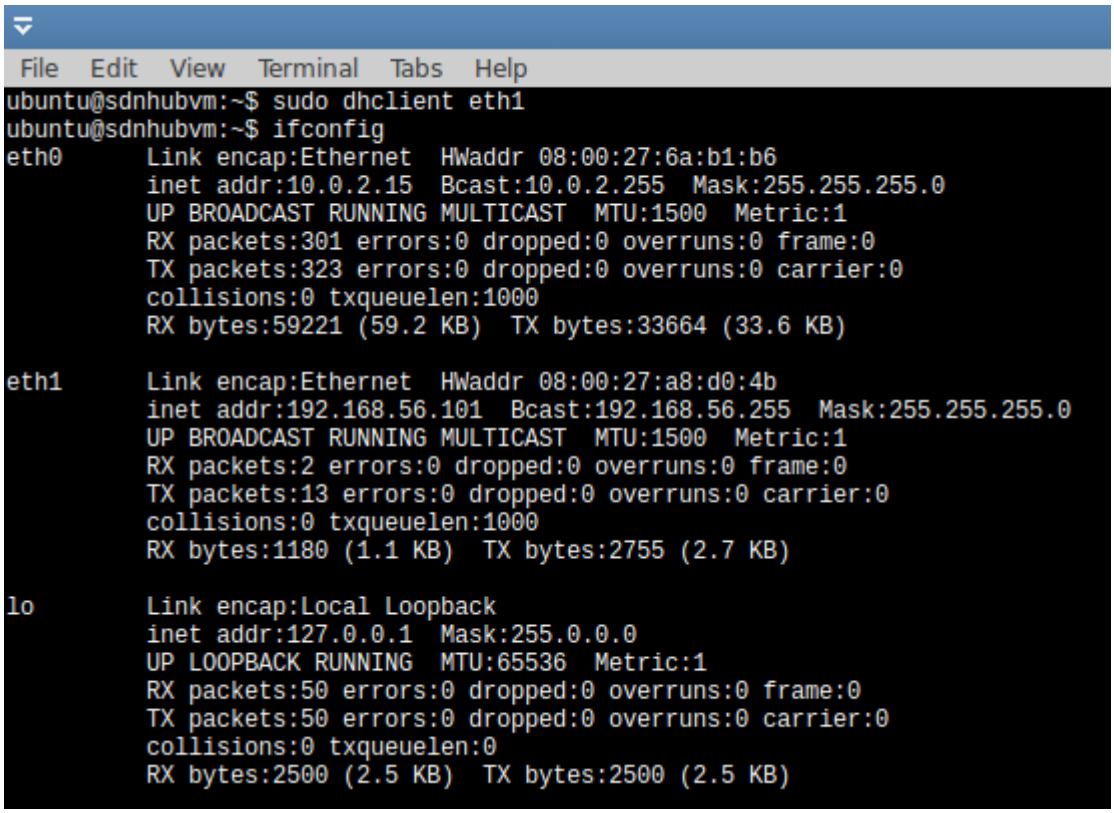

- Now, SSH into the virtual machine with the IP address of eth1 i.e 192.168..56.101
- This will now as for username and password, which in this case is Ubuntu in both the case.

- After successfully completing above steps you should now be able to work with the virtual machine.

2.2: Creating a mininet network -

- A network containing 3 hosts, 1 switch and 1 remote controller can be obtained by following command.
  Sudo mn –topo single,3 -- mac – switch ovsk – controller remote

```
ubuntu@sdnhubvm:~/pox$ sudo mn --topo single,3 --mac --switch ovsk --controller remote
*** Creating network
*** Adding controller
*** Adding hosts:
h1 h2 h3
*** Adding switches:
s1
*** Adding links:
(h1, s1) (h2, s1) (h3, s1)
*** Configuring hosts
h1 h2 h3
*** Starting controller
c0
*** Starting 1 switches
s1 ...
*** Starting CLI:
```

2.3: Wireshark packet capture and xterm -

- Open a new terminal and do the following to open wireshark packet capture.
  Sudo wireshark &
- Go to capture interface and select lo i.e loopback interface from the list.
- To open nodes h1 h2 h3 and to have graphical view of packet from on terminal put xterm h1 h2 h3 command.
- To see the traffic of packets flowing enter this command
  tcpdump –xx –n –I h-eth0
- To start the controller
  Cd pox
  ./pox.py log.level –DEBUG misc.of_tutorial
- For pinging between host and controller
  h1 ping –c1 h2

3) Software Design:

Basic commands which are used throughout the lab 1 is explained in the following section.

- **Sudo mn –** This includes one OpenFlow kernel switch connected to 2 hosts, plus the OpenFlow reference controller.
- **Sudo wireshark & -** To view traffic using the OpenFlow Dissector.
- **H1 ping –c1 h2 –** This will execute the ping command from host 0 to host 1 via ICMP echo request.
- **Xterm h1 h2 h3 –** This command will open all the nodes where graphical representation of packet flow can be seen.
- **Ifconfig –** With provide you with host address and Ip address.

Learning Switch Source Code:

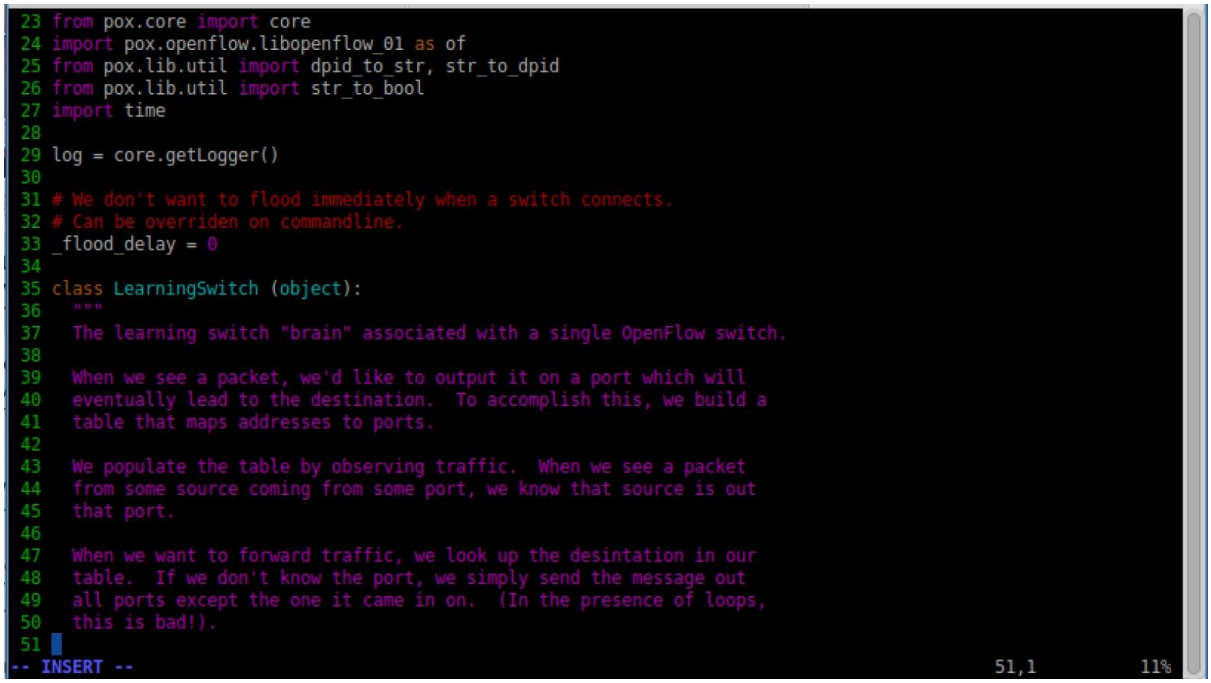

```
23 from pox.core import core
24 import pox.openflow.libopenflow_01 as of
25 from pox.lib.util import dpid_to_str, str_to_dpid
26 from pox.lib.util import str_to_bool
27 import time
28
29 log = core.getLogger()
30
31 # We don't want to flood immediately when a switch connects.
32 # Can be overriden on commandline.
33 _flood_delay = 0
34
35 class LearningSwitch (object):
36   """
37   The learning switch "brain" associated with a single OpenFlow switch.
38
39   When we see a packet, we'd like to output it on a port which will
40   eventually lead to the destination.  To accomplish this, we build a
41   table that maps addresses to ports.
42
43   We populate the table by observing traffic.  When we see a packet
44   from some source coming from some port, we know that source is out
45   that port.
46
47   When we want to forward traffic, we look up the desintation in our
48   table.  If we don't know the port, we simply send the message out
49   all ports except the one it came in on.  (In the presence of loops,
50   this is bad!).
51
-- INSERT --                                                    51,1          11%
```

```
  """
  def __init__ (self, connection, transparent):
    # Switch we'll be adding L2 learning switch capabilities to
    self.connection = connection
    self.transparent = transparent

    # Our table
    self.macToPort = {}

    # We want to hear PacketIn messages, so we listen
    # to the connection
    connection.addListeners(self)

    # We just use this to know when to log a helpful message
    self.hold_down_expired = _flood_delay == 0

    #log.debug("Initializing LearningSwitch, transparent=%s",
    #          str(self.transparent))

  def _handle_PacketIn (self, event):
    """
    Handle packet in messages from the switch to implement above algorithm.
    """

    packet = event.parsed

    def flood (message = None):
      """ Floods the packet """
      msg = of.ofp_packet_out()
-- INSERT --                                                    103,1         39%
```

```
      """ Floods the packet """
      msg = of.ofp_packet_out()
      if time.time() - self.connection.connect_time >= _flood_delay:
        # Only flood if we've been connected for a little while...

        if self.hold_down_expired is False:
          # Oh yes it is!
          self.hold_down_expired = True
          log.info("%s: Flood hold-down expired -- flooding",
              dpid_to_str(event.dpid))

        if message is not None: log.debug(message)
        #log.debug("%i: flood %s -> %s", event.dpid,packet.src,packet.dst)
        # OFPP_FLOOD is optional; on some switches you may need to change
        # this to OFPP_ALL.
        msg.actions.append(of.ofp_action_output(port = of.OFPP_FLOOD))
      else:
        pass
        #log.info("Holding down flood for %s", dpid_to_str(event.dpid))
      msg.data = event.ofp
      msg.in_port = event.port
      self.connection.send(msg)

    def drop (duration = None):
      """
      Drops this packet and optionally installs a flow to continue
      dropping similar ones for a while
      """
      if duration is not None:
-- INSERT --                                                    130,1         54%
```

```
129       """
130       if duration is not None:
131         if not isinstance(duration, tuple):
132           duration = (duration,duration)
133         msg = of.ofp_flow_mod()
134         msg.match = of.ofp_match.from_packet(packet)
135         msg.idle_timeout = duration[0]
136         msg.hard_timeout = duration[1]
137         msg.buffer_id = event.ofp.buffer_id
138         self.connection.send(msg)
139       elif event.ofp.buffer_id is not None:
140         msg = of.ofp_packet_out()
141         msg.buffer_id = event.ofp.buffer_id
142         msg.in_port = event.port
143         self.connection.send(msg)
144
145     self.macToPort[packet.src] = event.port # 1
146
147     if not self.transparent: # 2
148       if packet.type == packet.LLDP_TYPE or packet.dst.isBridgeFiltered():
149         drop() # 2a
150         return
151
152     if packet.dst.is_multicast:
153       flood() # 3a
154     else:
155       if packet.dst not in self.macToPort: # 4
156         flood("Port for %s unknown -- flooding" % (packet.dst,)) # 4a
157       else:
-- INSERT --                                                    157,1         68%
```

```
177 class l2_learning (object):
178   """
179   Waits for OpenFlow switches to connect and makes them learning switches.
180   """
181   def __init__ (self, transparent, ignore = None):
182     """
183     Initialize
184
185     See LearningSwitch for meaning of 'transparent'
186     'ignore' is an optional list/set of DPIDs to ignore
187     """
188     core.openflow.addListeners(self)
189     self.transparent = transparent
190     self.ignore = set(ignore) if ignore else ()
191
192   def _handle_ConnectionUp (self, event):
193     if event.dpid in self.ignore:
194       log.debug("Ignoring connection %s" % (event.connection,))
195       return
196     log.debug("Connection %s" % (event.connection,))
197     LearningSwitch(event.connection, self.transparent)
198
199
200 def launch (transparent=False, hold_down=_flood_delay, ignore = None):
201   """
202   Starts an L2 learning switch.
203   """
204   try:
205     global _flood_delay
-- INSERT --                                                    205,1         94%
```

```
201   """
202   Starts an L2 learning switch.
203   """
204   try:
205     global _flood_delay
206     _flood_delay = int(str(hold_down), 10)
207     assert _flood_delay >= 0
208   except:
209     raise RuntimeError("Expected hold-down to be a number")
210
211   if ignore:
212     ignore = ignore.replace(',', ' ').split()
213     ignore = set(str_to_dpid(dpid) for dpid in ignore)
214
215   core.registerNew(l2_learning, str_to_bool(transparent), ignore)
-- INSERT --                                                    215,1         Bot
```

Flowchart:

Following commands are assumed executed:

- Virtual Machine is running.
- Mininet Topology is created.
- Wireshark is configured.
- Nodes are opened and the ping command is given.
- Pox controller is initiated.

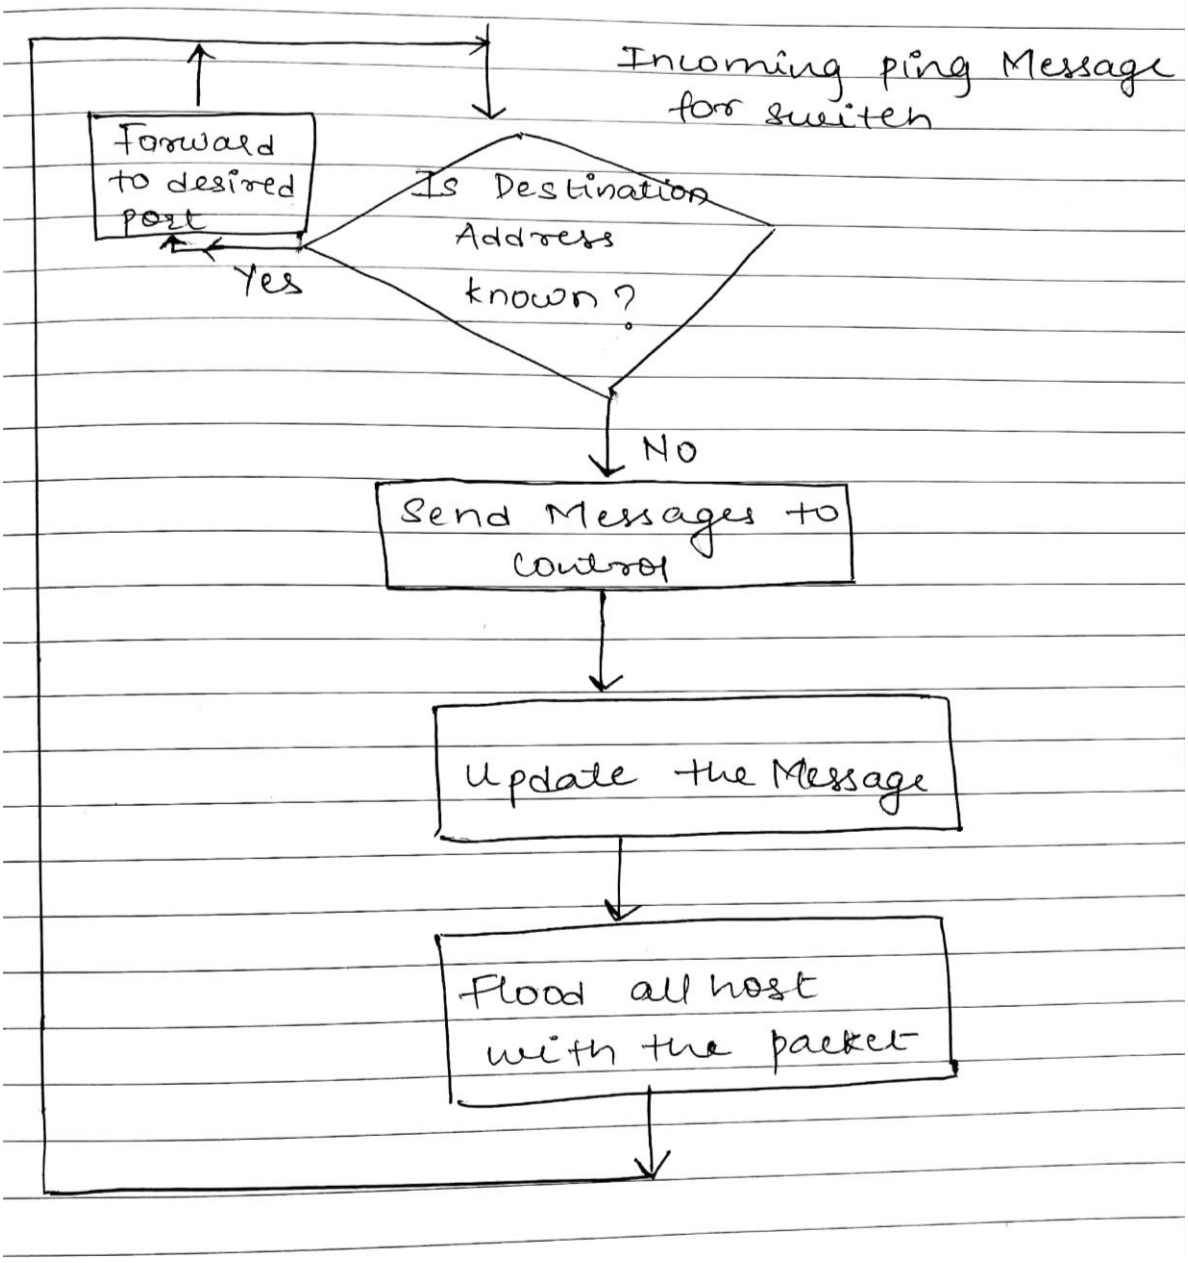

4) **Trouble shooting:**

- Understanding the concept of openflow and wireshark – The concepts and software were new so at the beginning it was difficult to understand why particular command worked in this way. To get used to the software I went through different videos and papers and understood the basic concept of wireshark, its use in packet flow and different commands that can be given to achieve the intended results.
- Connecting to the host controller – I was getting a message that controler is not connected to the system. This was due to SSH. I realised that my system name is sdnhubvm and not mininet. Difficult to SSH into the system was solved by first finding the IP address of eth1 and then SSH it using that IP address.
- Finding the IP Address of the eth1 – This was solved using sudo dhclient eth1 and ifconfig.

5) **Results:**

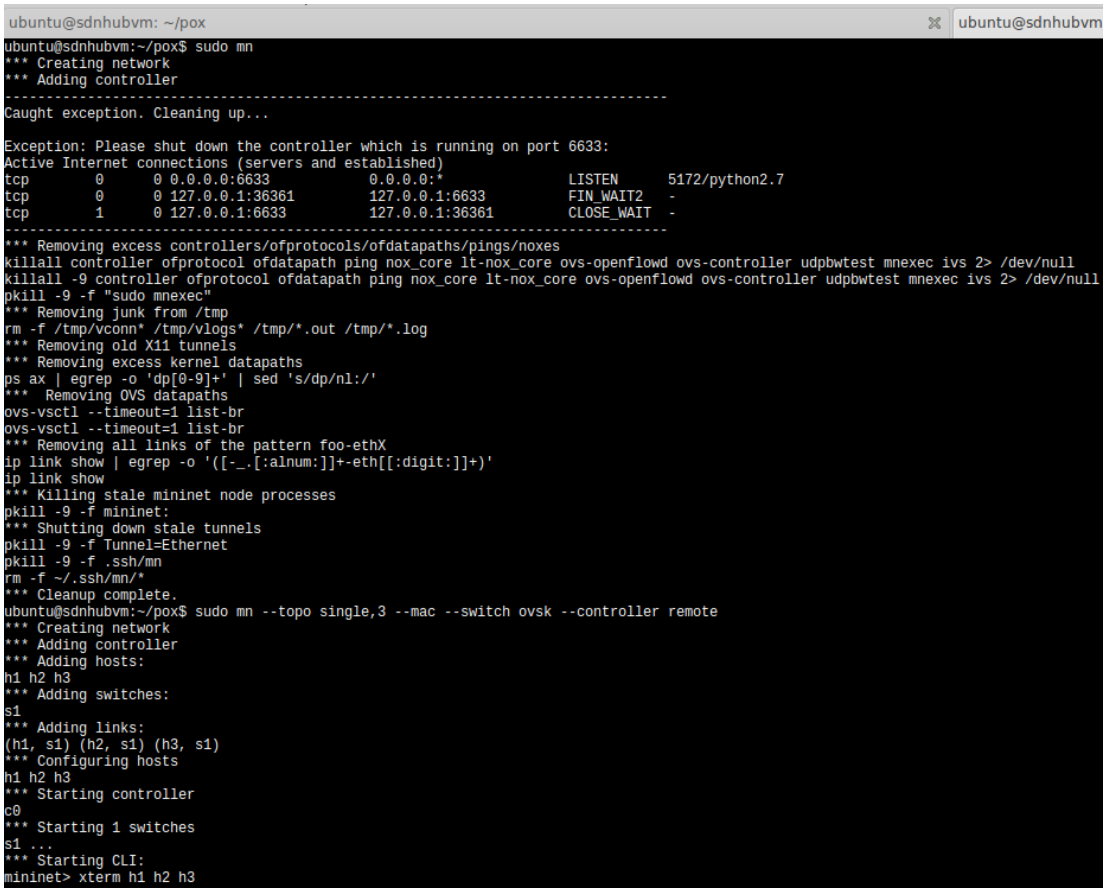

- The above image shows how the system makes a network comprising of host and controller and initiating links between them.

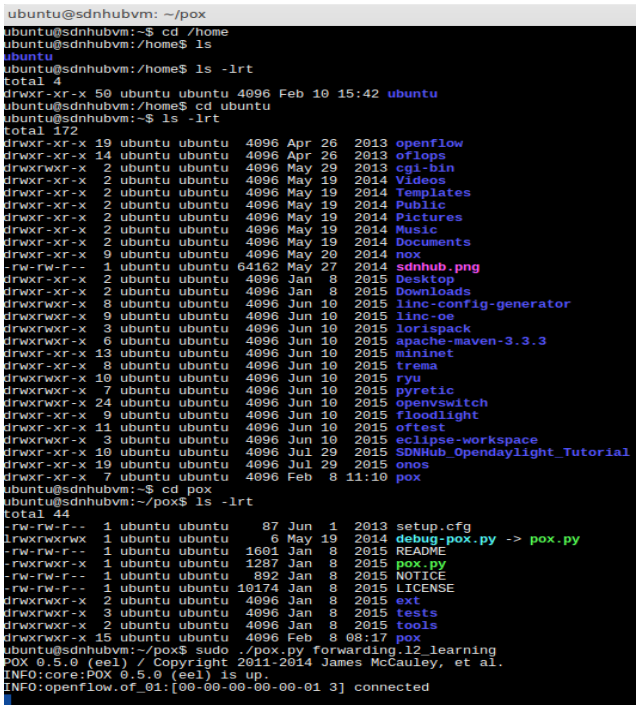

- Above images shows the directory that followed to reach the l2_learning python file which was used to study the switch. It is seen in the last line that openflow is connected.

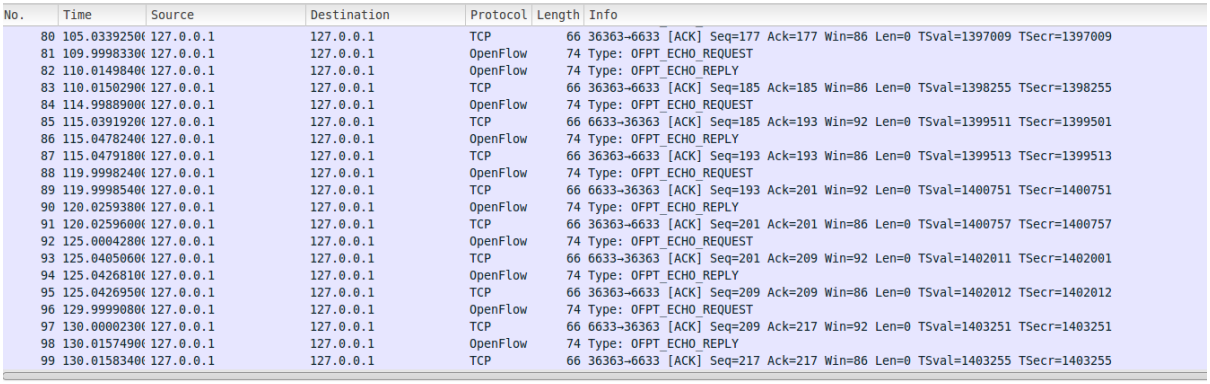

No.	Time	Source	Destination	Protocol	Length	Info
80	105.0339250(	127.0.0.1	127.0.0.1	TCP	66	36363→6633 [ACK] Seq=177 Ack=177 Win=86 Len=0 TSval=1397009 TSecr=1397009
81	109.9998330(	127.0.0.1	127.0.0.1	OpenFlow	74	Type: OFPT_ECHO_REQUEST
82	110.0149840(	127.0.0.1	127.0.0.1	OpenFlow	74	Type: OFPT_ECHO_REPLY
83	110.0150290(	127.0.0.1	127.0.0.1	TCP	66	36363→6633 [ACK] Seq=185 Ack=185 Win=86 Len=0 TSval=1398255 TSecr=1398255
84	114.9988900(	127.0.0.1	127.0.0.1	OpenFlow	74	Type: OFPT_ECHO_REQUEST
85	115.0391920(	127.0.0.1	127.0.0.1	TCP	66	6633→36363 [ACK] Seq=185 Ack=193 Win=92 Len=0 TSval=1399511 TSecr=1399501
86	115.0478240(	127.0.0.1	127.0.0.1	OpenFlow	74	Type: OFPT_ECHO_REPLY
87	115.0479180(	127.0.0.1	127.0.0.1	TCP	66	36363→6633 [ACK] Seq=193 Ack=193 Win=86 Len=0 TSval=1399513 TSecr=1399513
88	119.9998240(	127.0.0.1	127.0.0.1	OpenFlow	74	Type: OFPT_ECHO_REQUEST
89	119.9998540(	127.0.0.1	127.0.0.1	TCP	66	6633→36363 [ACK] Seq=193 Ack=201 Win=92 Len=0 TSval=1400751 TSecr=1400751
90	120.0259380(	127.0.0.1	127.0.0.1	OpenFlow	74	Type: OFPT_ECHO_REPLY
91	120.0259600(	127.0.0.1	127.0.0.1	TCP	66	36363→6633 [ACK] Seq=201 Ack=201 Win=86 Len=0 TSval=1400757 TSecr=1400757
92	125.0004280(	127.0.0.1	127.0.0.1	OpenFlow	74	Type: OFPT_ECHO_REQUEST
93	125.0405060(	127.0.0.1	127.0.0.1	TCP	66	6633→36363 [ACK] Seq=201 Ack=209 Win=92 Len=0 TSval=1402011 TSecr=1402001
94	125.0426810(	127.0.0.1	127.0.0.1	OpenFlow	74	Type: OFPT_ECHO_REPLY
95	125.0426950(	127.0.0.1	127.0.0.1	TCP	66	36363→6633 [ACK] Seq=209 Ack=209 Win=86 Len=0 TSval=1402012 TSecr=1402012
96	129.9999080(	127.0.0.1	127.0.0.1	OpenFlow	74	Type: OFPT_ECHO_REQUEST
97	130.0000230(	127.0.0.1	127.0.0.1	TCP	66	6633→36363 [ACK] Seq=209 Ack=217 Win=92 Len=0 TSval=1403251 TSecr=1403251
98	130.0157490(	127.0.0.1	127.0.0.1	OpenFlow	74	Type: OFPT_ECHO_REPLY
99	130.0158340(	127.0.0.1	127.0.0.1	TCP	66	36363→6633 [ACK] Seq=217 Ack=217 Win=86 Len=0 TSval=1403255 TSecr=1403255

- The above snapshot is from wireshark where open flow packets are traced real time. It follows OFPT protocol. Echo request and Echo reply Messages are exchanged between them.

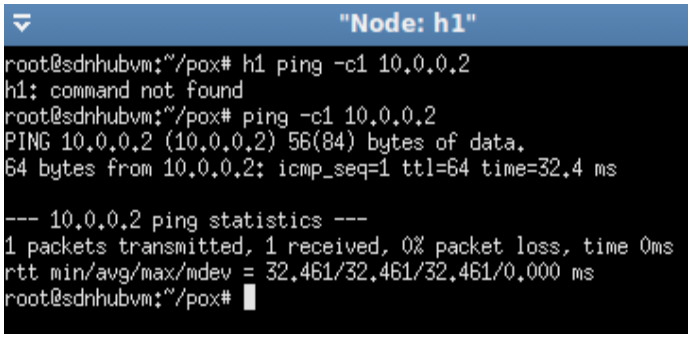

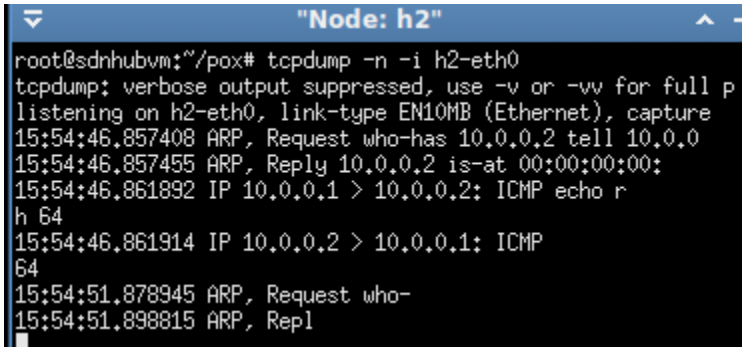

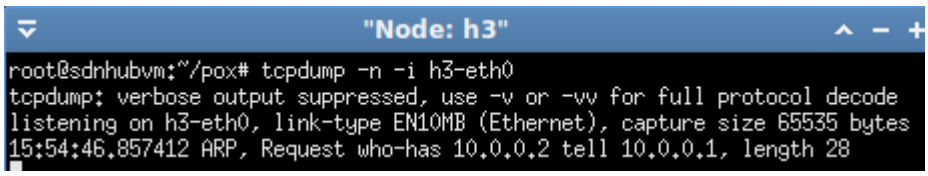

- The above three images shows the graphical representation of message flow.

6) Conclusion:

This lab thus made us understand the basic mininet terminology, for what it can be used. It also helped us learn the switch. When the openflow packets are passes through switch concepts and when packets are sent from controller to one respective host, all of them get the message just once. Once the controller Knows which host is it actually directed to, further packets are only sent to that host and thus flooding of packets at all ports is prevented.