

**SOFTWARE DEFINED NETWORKS**

**LAB 5**

**AKHILA NAIR**

**STUDENT ID – 01743358**

**HAND IN AND DUE DATE: 9<sup>th</sup> APRIL 2018**

**ESTIMATED HOURS OF WORKING: 10**

### 1. Purpose:

- To understand that single RX queue is not sufficient in high volume network.
- How to use Receive Side Scaling to increase the stability of DPDK.
- To understand the l3fwd example and ip-pipeline code and then inturn understand how to implement multiple core.
- To control each lcore to work on its own Rx queue.

### 2. Lab Setup:

- The setup remains the same as in other labs.
- We have to integrate the lab 2 source code to Implement receive side scaling and initiate multiple cores.

### 3. Software Design: Following are the changes made to the receiver.

- **Changing the Configuration of port:**

Here, to enable RSS, multimode is selected, and Ethernet RSS IPV4 is enabled.

```
.rxmode = {  
    .mq_mode = ETH_MQ_RX_RSS,  
    .header_split = 0, // Header split  
    .hw_ip_checksum = 0, // IP checksum  
offload  
    .hw_vlan_filter = 0, // VLAN filtering  
    .hw_vlan_strip = 0, // VLAN strip  
    .hw_vlan_extend = 0, // Extended VLAN  
    .jumbo_frame = 0, // Jumbo frame support  
    .hw_strip_crc = 0, // CRC strip by HW  
    .enable_scatter = 0, // Scattered packets  
RX handler  
    .max_rx_pkt_len = 9000, // Jumbo frame max  
packet len  
    .split_hdr_size = 0, // Header split  
buffer size  
},  
    .rx_adv_conf = {  
        .rss_conf = {  
            .rss_key = NULL,  
            .rss_key_len = 40,  
            .rss_hf = ETH_RSS_IPV4,  
        },  
    },  
}
```

- **Changing the number of RX queue:**

Since we need three Rx queue we have to set the number of rx queue which is a constant value to 3. This is done by using the API `rte_eth_dev_configure(pmd_id, No._of_rx_queue,0,&conf)`. `pmd_id` will always be zero.

```
status=rte_eth_dev_configure(pmd_id,3,0,&link_temp.conf);
```

- **Setting up the Mempool of each RX queue:**

Firstly, the mempool is created for each port,

```
charname1[128];sprintf(name1,"MEMPOOL1%u", pmd_id);
charname2[128];sprintf(name2,"MEMPOOL2%u", pmd_id);
charname3[128];sprintf(name3,"MEMPOOL3%u", pmd_id);
```

Then the pointer is assigned to each structure to point it out to the correct Queue.

```
struct rte_mempool * mp1;
struct rte_mempool * mp2;
struct rte_mempool * mp3;
```

The mempool name is changed to the current new name and other parameters remain same which is then mapped to `mp1`, `mp2` and `mp3` respectively.

```
mp1 = rte_mempool_create(
    name1, //mempoolname
    mempool_params_default.pool_size,
    mempool_params_default.buffer_size,
    mempool_params_default.cache_size,
    sizeof(struct rte_pktmbuf_pool_private),
    rte_pktmbuf_pool_init, NULL,
    rte_pktmbuf_init, NULL,
    socket_id,
    0}                                     );
```

```
if(mp1 == NULL) { printf("Error, can not create mempool for dev %u\n", pmd_id); exit(1); }
```

```

mp2 = rte_mempool_create(
    name2, //mempool name

    mempool_params_default.pool_size,
    mempool_params_default.buffer_size,
    mempool_params_default.cache_size,
    sizeof(struct rte_pktmbuf_pool_private),
    rte_pktmbuf_pool_init, NULL,
    rte_pktmbuf_init, NULL,
    socket_id,
    0);

```

```

if(mp2 == NULL) { printf("Error, can not create mempool for dev %u
\n", pmd_id); exit(1); }

```

```

mp3 = rte_mempool_create(
    name3, //mempool name
    mempool_params_default.pool_size,
    mempool_params_default.buffer_size,
    mempool_params_default.cache_size,
    sizeof(struct rte_pktmbuf_pool_private),
    rte_pktmbuf_pool_init, NULL,
    rte_pktmbuf_init, NULL,
    socket_id,
    0);

```

```

if(mp3 == NULL) { printf("Error, can not create mempool for dev %u
\n", pmd_id); exit(1); }

```

- **Configuring the RX queue:**

```

status = rte_eth_rx_queue_setup(
                                pmd_id,
                                0,
                                default_hwq_in_params.size,
                                socket_id,
                                &default_hwq_in_params.conf,
                                mp1 //memory pool address
                                );

```

```

if (status < 0) { printf("Error, can not set up queue for dev %u \n",
pmd_id); exit(1); }

```

```

        status = rte_eth_rx_queue_setup(
                                pmd_id,
                                1,
                                default_hwq_in_params.size,
                                socket_id,
                                &default_hwq_in_params.conf,
                                mp2 //memory pool address
                                );

if (status < 0) { printf("Error, can not set up queue for dev %u \n",
pmd_id); exit(1); }

```

```

        status = rte_eth_rx_queue_setup(
                                pmd_id,
                                2,
                                default_hwq_in_params.size,
                                socket_id,
                                &default_hwq_in_params.conf,
                                mp3 //memory pool address
                                );

if (status < 0) { printf("Error, can not set up queue for dev %u \n",
pmd_id); exit(1); }

```

- **Setting up each core to work on a different RX queue:**  
For each lcore\_id in the burst function queue ID will be changed respectively.

```

        if(lcore_id == 0)
        {
            printf("Hello from master core 0 !\n",
lcore_id);

            while(1)
            {
                uint32_t total_time_in_sec = 10;
                uint64_t p_ticks = total_time_in_sec * rte_get_tsc_hz();
                uint64_t p_start = rte_get_tsc_cycles();
                uint32_t total_pkts = 0; //for statistics
                while(rte_get_tsc_cycles() - p_start < p_ticks)
                {
                    n_pkts=rte_eth_rx_burst(0,0,pkts,RTE_PORT_IN_BURST_SIZ
E_MAX); //trying to receive packets
                    if(unlikely(n_pkts == 0)) {continue;} //if no packet
received, then start the next try
                }
            }
        }

```

```

        total_pkts += n_pkts;
        //retrieving the data from each packet
    }
    else if(lcore_id == 1)
    {
        printf("Hello from master core 1 !\n",
lcore_id);

        while(1)
        {
uint32_t total_time_in_sec = 10;
uint64_t p_ticks = total_time_in_sec * rte_get_tsc_hz();
uint64_t p_start = rte_get_tsc_cycles();
uint32_t total_pkts = 0; //for statistics
while(rte_get_tsc_cycles() - p_start < p_ticks)
        {
n_pkts=rte_eth_rx_burst(0,1,pkts,RTE_PORT_IN_BURST_SIZ
E_MAX); //trying to receive packets
if(unlikely(n_pkts == 0)) {continue;} //if no packet
received, then start the next try
total_pkts += n_pkts;

else if(lcore_id == 2)
        {
            printf("Hello from master core 2 !\n",
lcore_id);

            while(1)
            {
uint32_t total_time_in_sec = 10;
uint64_t p_ticks=total_time_in_sec* rte_get_tsc_hz();
uint64_t p_start = rte_get_tsc_cycles();
uint32_t total_pkts = 0; //for statistics
while(rte_get_tsc_cycles() - p_start < p_ticks)
            {
n_pkts=rte_eth_rx_burst(0,2,pkts,RTE_PORT_IN_BURST_SIZ
E_MAX);
if(unlikely(n_pkts == 0)) {continue;}
total_pkts += n_pkts;
//retrieving the data from each packet

```

#### 4. Troubleshooting:

- At a point of time, wrong build file was getting created.  
I discussed it with TA and did the necessary changes in the make file and the issue was solved.
- Other issue was that whenever the lcore calls the `rte_eth_rx_burst()`, each lcore was passing the same queue ID and hence received packets only in first core.  
After thoroughly understanding the function of `rte_eth_rx_burst()`, and finding out the parameters used here, different queue ID was passed so that each lcore will handle its own RX queue.
- I got “expected identifier or ‘(‘ before ‘)’ token” error since the main function was not closed correctly. This issue was resolved by carefully going through the code and doing the necessary changes.

#### 5. Results:

- The following image shows the working of TX side and the command used to run the program is

```
sudo ./build/tx_demo -w 0000:03:00.0 --socket-mem 64 --file-prefix tx -c 0x08
```

```
^C[test@localhost TX]make
/home/test/dpdk/lab2/TX/Makefile:19: warning: overriding recipe for target `clean'
/usr/share/dpdk/mk/rte.app.mk:270: warning: ignoring old recipe for target `clean'
[test@localhost TX]$ sudo ./build/tx_demo -w 0000:03:00.0 --socket-mem 64 --file-prefix tx -c 0x08
[sudo] password for test:
EAL: Detected 4 lcore(s)
EAL: Probing VFIO support...
EAL: PCI device 0000:03:00.0 on NUMA socket -1
EAL:   probe driver: 8086:1533 net_e1000_igb
Getting the MAC address of the selected port...
00:08:a2:0c:03:a3:
=====Setting up the RXQ0.0=====
=====Setting up the TXQ0.3=====
PMD: eth_igb tx_queue_setup(): To improve 1G driver performance, consider setting the TX WTHRESH value to 4, 8, or 16.
Hello from core 3!!!
lcore 3: Throughput is 0.453745 GBPS
lcore 3: Throughput is 0.727279 GBPS
lcore 3: Throughput is 0.516647 GBPS
lcore 3: Throughput is 0.727279 GBPS
lcore 3: Throughput is 0.727279 GBPS
```

- The following image shows the working of RX side and the command used to run the program is

```
sudo ./build/rx_demo -c 0x07 -w 0000:02:00.0 --socket-  
mem 256 --file-prefix rx -
```

```
test@localhost RX]$ make  
/home/test/dpdk/lab2/RX/Makefile:19: warning: overriding recipe for target `clean'  
/usr/share/dpdk/mk/rte.app.mk:270: warning: ignoring old recipe for target `clean'  
CC main.o  
LD rx_demo  
INSTALL-APP rx_demo  
INSTALL-MAP rx_demo.map  
test@localhost RX]$ sudo ./build/rx_demo -c 0x07 -w 0000:02:00.0 --socket-mem 256 --file-prefix rx --  
[sudo] password for test:  
EAL: Detected 4 lcore(s)  
EAL: Probing VFIO support...  
EAL: PCI device 0000:02:00.0 on NUMA socket -1  
EAL: probe driver: 8086:1533 net_e1000_igb  
##### Setting UP RXQ For PORT 0 #####  
Hello from master core 1 !  
Hello from master core 2 !  
Hello from master core 0 !  
core 1, received 2793246 packets in 10 seconds.  
core 2, received 2793438 packets in 10 seconds.  
core 0, received 3581454 packets in 10 seconds.  
core 1, received 4327949 packets in 10 seconds.  
core 2, received 4327953 packets in 10 seconds.  
core 0, received 5548702 packets in 10 seconds.
```

## 6. Conclusion:

Hence multiple lcore and thus queue was enabled for high volume network. Used Receive Side Scaling efficiently to increase the stability of DPDK.